Oct 31 — 5 papers

Episode 43 · October 31, 2026 · 5 papers

Safe Platooning under Input Constraints with Link-Wise Velocity String Stability

SKIM

Panagiotis S. Trakas, Charalampos P. Bechlioulis, Thomas Berger

A funnel-control platoon controller that uses only onboard gap and speed sensing, respects each car's own braking and throttle limits, and comes with proofs of no collisions and bounded speed deviation that do not grow with platoon length.

PROBLEM Most platoon controllers with string-stability guarantees either ignore actuator saturation, assume identical vehicles, or depend on online optimization like MPC. Guarantees on safety, input limits and string stability together for mixed-capability vehicles that can stop and restart have been hard to get without exact models.

METHOD Each follower runs a model-free funnel controller on its spacing error. The spacing target includes a braking-distance margin that accounts for different deceleration limits, and the funnel adaptively widens when propulsion saturates so the controller does not demand thrust it cannot deliver. A standstill rule stops vehicles from rolling backward, and the proofs cover existence of solutions, collision avoidance, input-limit satisfaction, and a per-link velocity bound with gain one plus an explicit offset that does not depend on vehicle index.

RESULTS The evidence is mainly the theorems. The abstract gives no numbers beyond simulations with 2 to 100 followers under disturbances and actuator limits, and it does not mention a baseline or any measured error, spacing, or comparison with CACC or MPC.

CAVEAT Gain one with an additive offset is a weaker notion than standard string stability, because deviations are not attenuated along the string, and the offset scales with the funnel width and braking margin, so tighter guarantees probably mean larger gaps and lower throughput. The guarantees also rest on explicit uniform conditions on parameters and initial states, the controller probably still needs bounds on the predecessor's braking capability, and validation is simulation only with no communication at all, so there is nothing to say about packet loss or delay.

WHY IT MATTERS This is peripheral to your work. It is a sensor-only, predecessor-following ACC design with no message exchange, no joining or splitting maneuvers, and no negotiation, so it does not touch intent sharing or protocol design. It could serve as a provably safe fallback layer for when V2V links drop, or as a non-communicating baseline against which to quantify what CACC buys you.

arXiv:2610.39964v1 PDF incremental

A Bi-level Co-design Framework for Lunar Monoped with 2D Jump and Landing Control

SKIM

Sanjay Ramkumar, Suraj Kumar

They tune the leg lengths of a single-legged hopping robot for moon gravity and show that the best leg depends on which jump controller you pair it with.

PROBLEM In lunar gravity, contact forces are weak and flight phases are long, so a leg's shape and its controller depend on each other. Designing them separately can give a mismatched robot.

METHOD An outer loop uses differential evolution to search over the robot's two link lengths. Each candidate leg is scored on a vertical jump using one of two inner controllers, a quadratic-program task-space controller or a direct-collocation trajectory optimizer. For two-dimensional jumps, they add a trajectory optimizer that regulates centroidal angular momentum and an LQR controller that absorbs the landing impact.

RESULTS Trajectory optimization beats the task-space controller across most of the design space, with a seventeen percent lower cost of transport when each controller is compared at its own best design. The two-dimensional jumps, both in place and forward, are shown only on one fixed design with two links of zero point three meters each, and the abstract gives no numbers for them.

CAVEAT It is all planar simulation with no hardware and no terrain variation. The co-design search covers only vertical jumps and is never applied to the two-dimensional jumps. The seventeen percent gap mostly shows the expected advantage of offline trajectory optimization over a reactive controller, which is not a fair comparison for real-time use.

WHY IT MATTERS This is outside your area. It is single-robot legged locomotion with no communication, no multiple agents, and no negotiation, so nothing carries over to maneuver coordination, platooning, or conflict resolution. The only loose link is the idea of designing a system together with its controller, and that is not specific to this paper.

arXiv:2610.37199v1 PDF incremental

GAWM: A Compact Latent World-Action Model for Multi-Task Robot Manipulation

SKIM

Yang Liu et al.

A model under a hundred million parameters that predicts future DINOv3 features instead of future pixels, then reads robot actions off those predictions, and it matches or beats billion-parameter manipulation models on RoboTwin and LIBERO.

PROBLEM World-action models for robot manipulation usually sit on billion-parameter vision-language or video-generation backbones, so they are expensive to train and slow at inference. It was unclear whether the benefit of predicting the future survives without that scale and without robot or video pretraining.

METHOD Each camera view goes through a frozen DINOv3 encoder and is compressed into a few tokens. A small predictor, conditioned on the task and not on actions, forecasts those tokens at two future horizons. That forecast is supervised against the dense patch features of the real future frames, which forces it to encode spatial layout. An action decoder then regresses a chunk of future actions directly from the current and predicted tokens.

RESULTS The abstract reports no numbers. It claims GAWM beats billion-parameter world-action models on RoboTwin 2.0 and roughly matches billion-parameter VLAs on LIBERO. On a real dual-arm robot it gets lower action-prediction error than a fine-tuned billion-parameter VLA on held-out demonstrations, with much lower latency.

CAVEAT The real-robot result is offline action-prediction error on held-out demonstrations, not closed-loop task success, so it says little about deployed performance. The ablation shows the world model helps mainly when data or training time is limited, which suggests the gain shrinks at full scale. Comparable to on LIBERO also means it does not beat VLAs there.

WHY IT MATTERS This is peripheral to your work. It is single-agent robot manipulation with no communication, no negotiation and no multiple vehicles. The only transferable idea is that a small model predicting future latent features from a frozen encoder can stand in for a large generative backbone, which might matter if you ever need compact on-board intent prediction. It does not touch maneuver coordination or message design.

arXiv:2610.37091v1 PDF incremental

Polarity--Time Interaction and Multi-Level Reliability for Event-Based Spacecraft Pose Estimation

SKIM

Yonglong Zhang, Zongwu Xie, Yang Liu

An event-camera pipeline for estimating spacecraft pose that weights each keypoint and each frame by how reliable it looks, and edges out the previous best entry on the SPADES challenge.

PROBLEM With event cameras, sparse and uneven event responses make some keypoints much less accurate than others. Standard pose pipelines treat every correspondence and every frame as equally trustworthy, so a few bad keypoints or one poor frame can corrupt the pose and the trajectory built from it.

METHOD Positive events, negative events and a map of recent-event timing each go through their own branch, and the three branches are fused with cross-branch attention and gating. The keypoint head predicts a precision value for each correspondence, and those values weight a geometrically constrained pose solver. Two refinement steps follow. The first aligns detected event edges with the projected CAD model. The second smooths the trajectory over the sequence, scaling the trust in each frame with a hand-built ten-feature reliability descriptor covering heatmap shape, reprojection consistency and geometric conditioning.

RESULTS On the official SPADES test set the full system scores 0.0474, about 5.4 percent lower than PASAN plus IRIS, and lower is better on this metric. In validation ablations, the complete cascade improves the score from 0.0772 for plain localization with uniform-weight PnP to 0.0514. On identical single-frame inputs, quality weighting accounts for only part of that gain, moving from 0.0549 to 0.0514.

CAVEAT Everything is tuned and tested on a single challenge benchmark. The margin over the best prior entry is small, and the paper gives no variance or significance test to show it is real. Most of the improvement comes from stacking several refinement stages, so the headline idea of reliability weighting buys only about seven percent on its own. The ten-feature descriptor is handcrafted and probably specific to this dataset. The abstract also gives no runtime figures, which matters for onboard use.

WHY IT MATTERS This paper is outside your area. It covers monocular pose estimation for one spacecraft and involves no communication, negotiation or multiple agents. The only idea you could borrow is general: weight each observation by an estimated reliability before trajectory optimization, similar to discounting stale or low-quality V2X intent messages. You can get that idea from the abstract without reading the paper.

arXiv:2610.36625v1 PDF incremental

RadarCom: Unlocking Simultaneous Communication and Sensing on Commodity Automotive mmWave Radars SKIM

Yi Wei et al.

They turn the interference between two stock TI automotive radars into a few-kilobit-per-second data link, with no hardware changes, and the radars keep sensing while it runs.

PROBLEM V2V messaging needs dedicated radios that most cars still lack, while nearly every production car already carries a mmWave radar. Mutual interference between those radars has always been treated as noise to suppress, so nobody had tried to use it as a channel on unmodified commercial hardware.

METHOD A transmitting radar encodes data in its chirps, and the receiving radar reads that data out of the interference it would normally throw away. Each radar finds the other by noticing a mismatch in chirp slope, then locks its chirp timing to the other radar's. Symbols are differentially encoded on the range-FFT, and an interference suppression step tuned for sensing keeps the radar's own detections usable. Everything runs in software on TI AWR2243 boards, with no change to the waveform or the hardware.

RESULTS On AWR2243 hardware in static, dynamic and real-road tests, they report an average effective rate of 5.55 kilobits per second, symbol error rates below 1.5 percent, and a sensing F1 score of at most 80.5 percent during communication. The abstract gives no figure for sensing without communication and no comparison against any baseline.

CAVEAT The headline sensing number is a best case, and without a no-communication baseline you cannot tell how much detection quality the link costs. The abstract says nothing about range, how many vehicles can share the link, latency, or what happens when the radars are not pointed at each other. That last point matters because a forward-facing radar mostly sees the back of the car ahead, not a car beside it negotiating a merge.

WHY IT MATTERS This is peripheral to your work. It is about the physical link, and it says nothing about negotiation protocols, message design or planning. One angle is still worth keeping in mind. A low-rate, line-of-sight radar link between a leader and its follower could serve as a fallback channel for platoon or CACC messages. You could then add it as one more degraded-communication scenario when you test how coordination holds up with lossy or partial links. A few kilobits per second is enough for short intent or acceptance messages, but not for full trajectory sharing.

arXiv:2610.36622v1 PDF notable